

AWS Spark Wine Quality Prediction Application

pySpark AWS Wine Prediction App Is an application that was developed with Python's pySpark interface and installed on the AWR EMR cluster.

The purpose of this project is to leverage publicly available data to train a machine learning model in parallel on EC2 instances, with the goal of predicting wine quality. In order to make deployments easier, the project also uses Docker to produce a container image for a trained machine learning model.

This project contains 2 main python source files:

wine_prediction.py reads training dataset from S3 and trains model in parallel on EMR spark cluster. Once model is trained, it can be run on provided test data provided via S3. This program stores trained model in S3 bucket (Public URL for bucket - S3://pa2winequ)

wine_prediction.py program loads trained model and executes that model on given test data file. This will then print F1 score as metrics for accuracy of the trained model.

Docker file: Docker file to create docker image and run container for simplified deployment.

Link to GitHub code

<https://github.com/vijayasindhudandeboina/AWS-Spark-wine-Quality>

Link to docker container image

<https://hub.docker.com/repository/docker/vijaya430/wine-prediction-app/general>

Source files:

`wine_prediction.py` reads Training dataset from S3 and trains model in parallel on EMR spark cluster. Once model is trained, it can be run on provided test data provided via S3. This program stores trained model in S3 bucket (Public URL for bucket - S3://wine-data-12)

`wine_test_data_prediction.py` program loads trained model and executes that model on given test data file. This will then print F1 score as metrics for accuracy of the trained model.

`Docker file` creates docker image and run container for simplified deployment.

Instruction to use:

Phase 1: Pre-requisites

Prepare the PEM file, credentials, and AWS account (same as last assignment).

Phase 2: Copy data to S3.

We will copy all the datasets given and the python scripts to s3 in our bucket.

Phase 3: Establish an EMR Cluster

1. Create Spark cluster in AWS.

Using the EMR console that AWS provides, users can establish spark clusters.

Kindly follow the instructions to build one using four EC2 instances (you can use more, depending on your load).

1. Use the "EC2-> Network & Security -> Key-pairs" navigation to create a key pair for the EMR cluster.

Utilize the format pem. The {name of key pair}>.pem file will be downloaded as a result. To perform an SSH to an EC2 instance, you will need to keep it secure.

2. Go to the console for Amazon EMR. Next, select Clusters -> Create Cluster from the menu.

3. Now complete the corresponding sections:

Proceed to the General Configuration -> Cluster Name -> Software Configuration -> EMR 5.33 menu and choose 'Spark: Spark 2.4.7 on Hadoop 2.10.1 YARN and Zeppelin 0.9.0'.

Hardware Configuration -> Set the number of instances to 4.

Pem key generated in the previous step under Security Access -> Provide.

The remaining settings can be left as-is.

4. Following a successful cluster creation, the cluster status should be "Waiting."

us-east-1.console.aws.amazon.com/emr/home?region=us-east-1#/doneCluster/j-VEO4DB2K6JFI

Clone "pa2_wine_qu" Info

Name and applications Info

Name
pa2_wine_qu

Amazon EMR release Info
A release contains a set of applications which can be installed on your cluster.
emr-5.33.0

Application bundle

Spark

Core Hadoop

HBase

Presto

Custom

☐ Flink 1.12.1
☐ HCatalog 2.3.7
☐ Hue 4.9.0
☐ Livy 0.7.0
☐ Oozie 5.2.0
☐ Presto 0.245.1
☐ TensorFlow 2.4.1
☐ ZooKeeper 3.4.14

☐ Ganglia 3.7.2
☒ Hadoop 2.10.1
☒ JupyterEnterpriseGateway 2.1.0
☐ MXNet 1.7.0
☐ Phoenix 4.14.3
☒ Spark 2.4.7
☐ Tez 0.9.2

☐ HBase 1.4.13
☐ Hive 2.3.7
☐ JupyterHub 1.1.0
☐ Mahout 0.13.0
☐ Pig 0.17.0
☐ Squirrel 1.4.7
☒ Zeppelin 0.9.0

Summary Info

Name and applications

Name
pa2_wine_qu

Amazon EMR release
emr-5.33.0

Application bundle
Custom (Hadoop 2.10.1, JupyterEnterpriseGateway 2.1.0, Spark 2.4.7, Zeppelin 0.9.0)

Cluster configuration

Instance groups
Primary (m5.xlarge), Core (m5.xlarge), Task (m5.xlarge)

Cluster scaling and provisioning

Provisioning configuration
Core size: 1 instance

Activate Windows
Go to Settings to activate Windows.

us-east-1.console.aws.amazon.com/emr/home?region=us-east-1#/doneCluster/j-VEO4DB2K6JFI

AWS Glue Data Catalog settings

Use the AWS Glue Data Catalog to provide an external metastore for your application.

☐ Use for Spark table metadata

Custom Amazon Machine Image (AMI) Info

Choose or enter an AMI ID

☒ Update all installed packages on reboot

Cluster configuration Info

Choose a configuration method for the primary, core, and task node groups for your cluster.

☒ Instance groups
Choose one instance type per node group

☐ Instance fleets
Choose any combination of instance types within each node group

Instance groups

Primary

Choose EC2 instance type

m5.xlarge
4 vCore 16 GiB memory EBS only storage
On-Demand price: - Lowest Spot price: -

Actions

Summary Info

Name and applications

Name
pa2_wine_qu

Amazon EMR release
emr-5.33.0

Application bundle
Custom (Hadoop 2.10.1, JupyterEnterpriseGateway 2.1.0, Spark 2.4.7, Zeppelin 0.9.0)

Cluster configuration

Instance groups
Primary (m5.xlarge), Core (m5.xlarge), Task (m5.xlarge)

Cluster scaling and provisioning

us-east-1.console.aws.amazon.com/emr/home?region=us-east-1#/doneCluster/fj-VEO4DB2K6UFI

Services Search [Alt+S]

Core

Choose EC2 instance type

m5.xlarge
4 vCore 16 GiB memory EBS only storage
On-Demand price: - Lowest Spot price: -

Actions ▾

► Node configuration - optional

Task 1 of 1

Name

Task - 1

Choose EC2 instance type

m5.xlarge
4 vCore 16 GiB memory EBS only storage
On-Demand price: - Lowest Spot price: -

Actions ▾

► Node configuration - optional

Add task instance group

You can add up to 47 more task instance groups.

Remove instance group

Summary Info

Name and applications

Name
pa2_wine_qu

Amazon EMR release
emr-5.33.0

Application bundle
Custom (Hadoop 2.10.1,
JupyterEnterpriseGateway 2.1.0, Spark 2.4.7,
Zeppelin 0.9.0)

Cluster configuration

Instance groups
Primary (m5.xlarge), Core (m5.xlarge), Task (m5.xlarge)

Cluster scaling and provisioning

Provisioning configuration
Core size: 1 instance

us-east-1.console.aws.amazon.com/emr/home?region=us-east-1#/doneCluster/fj-VEO4DB2K6UFI

Services Search [Alt+S]

Cluster scaling and provisioning Info

Set up scaling and provisioning configurations for the core and task node groups for your cluster.

Choose an option

☒ Set cluster size manually
Use this option if you know your workload patterns in advance.

☐ Use EMR-managed scaling
Monitor key workload metrics so that EMR can optimize the cluster size and resource utilization.

☐ Use custom automatic scaling
To programmatically scale core and task nodes, create custom automatic scaling policies.

Provisioning configuration

Set the size of your core and task instance groups. Amazon EMR attempts to provision this capacity when you launch your cluster.

Name	Instance type	Instance(s) size	Use Spot purchasing option
Task - 1	m5.xlarge	3	<input type="checkbox"/>
Core	m5.xlarge	1	<input type="checkbox"/>

Networking Info

Virtual private cloud (VPC) Info

vpc-0e74ca151a6acd930

Browse Create VPC

Subnet Info

Summary Info

Core node security group
sg-06e566b184...

Cluster termination

Cluster termination
Manually terminate cluster

Security configuration and EC2 key pair - optional

Amazon EC2 key pair
pa2_wine_qu

Identity and Access Management (IAM) roles

Service role
EMR_DefaultRole

Instance profile
EMR_EC2_DefaultRole

us-east-1.console.aws.amazon.com/emr/home?region=us-east-1#/cloneCluster/fj-VE04DB2K6UFI

Subnet [Info](#)

subnet-00d2ac4156bd44155 [Browse](#) [Create subnet](#)

EC2 security groups (firewall)

Steps - optional (0) [Info](#) [Remove](#) [Edit](#) [Add](#)

Use commands and scripts to tell your cluster where to find and how to process your data. Steps run consecutively unless you enable the Concurrency option.

Cluster termination [Info](#)

☒ Manually terminate cluster

☐ Automatically terminate cluster after last step ends

☐ Automatically terminate cluster after idle time (Recommended)

☐ Use termination protection

Protect your EC2 instances from accidental termination.

Bootstrap actions - optional [Info](#)

Use bootstrap actions to install software or customize your instance configuration.

Cluster logs - optional [Info](#)

Summary [Info](#)

Core node security group
[sg-06e566b184...](#)

Cluster termination

Cluster termination
Manually terminate cluster

Security configuration and EC2 key pair - optional

Amazon EC2 key pair
[pa2_wine_qu](#)

Identity and Access Management (IAM) roles

Service role
[EMR_DefaultRole](#)

Instance profile
[EMR_EC2_DefaultRole](#)

Activate Windows
Go to Settings to activate Windows.

us-east-1.console.aws.amazon.com/emr/home?region=us-east-1#/cloneCluster/fj-VE04DB2K6UFI

Security configuration and EC2 key pair - optional [Info](#)

Security configuration

Select your cluster encryption, authentication, authorization, and instance metadata service settings.

[Choose a security configuration](#) [Browse](#) [Create security configuration](#)

Amazon EC2 key pair for SSH to the cluster [Info](#)

[pa2_wine_qu](#) [Browse](#) [Create key pair](#)

Identity and Access Management (IAM) roles [Info](#)

Choose or create a service role and instance profile for the EC2 instances in your cluster.

Amazon EMR service role [Info](#)

The service role is an IAM role that Amazon EMR assumes to provision resources and perform service-level actions with other AWS services.

☒ Choose an existing service role

Select a default service role or a custom role with IAM policies attached so that your cluster can interact with other AWS services.

☐ Create a service role

Let Amazon EMR create a new service role so that you can grant and restrict access to resources in other AWS services.

Service role

[EMR_DefaultRole](#)

Summary [Info](#)

Core node security group
[sg-06e566b184...](#)

Cluster termination

Cluster termination
Manually terminate cluster

Security configuration and EC2 key pair - optional

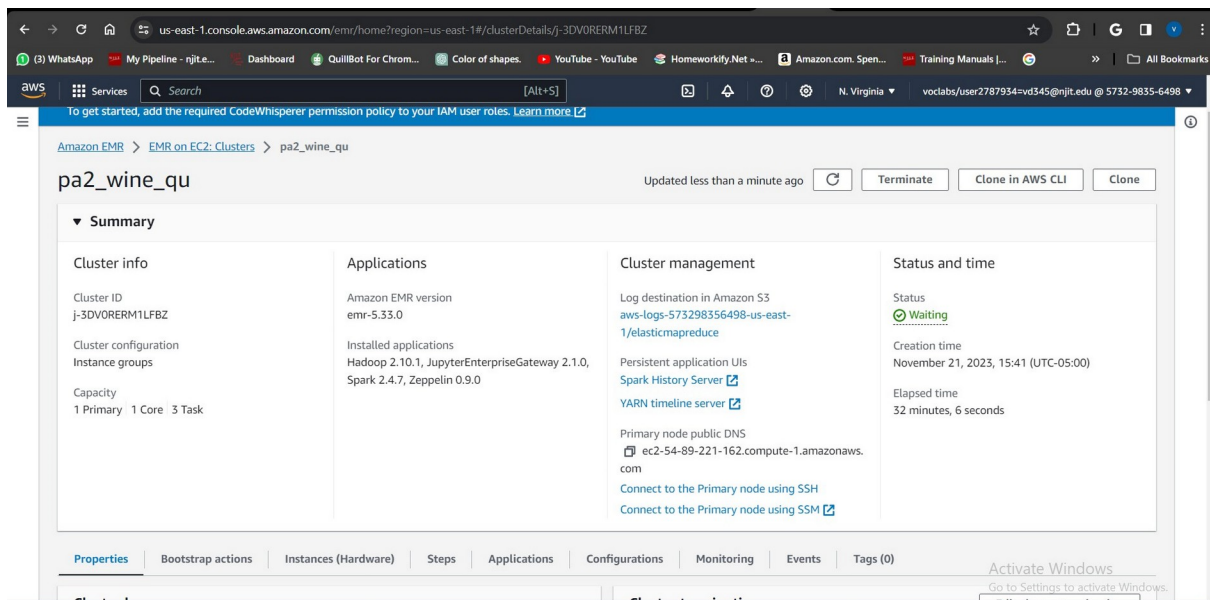
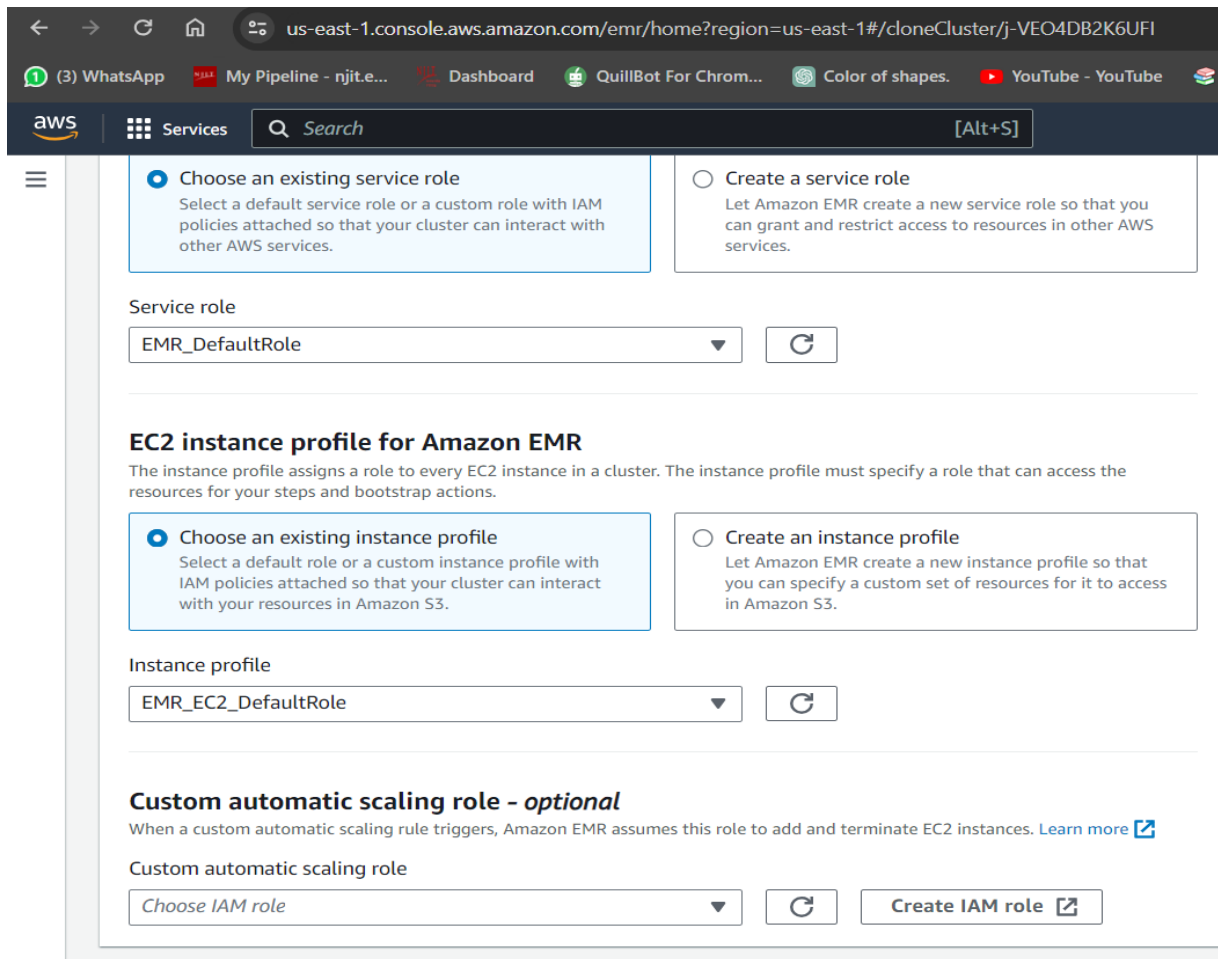
Amazon EC2 key pair
[pa2_wine_qu](#)

Identity and Access Management (IAM) roles

Service role
[EMR_DefaultRole](#)

Instance profile
[EMR_EC2_DefaultRole](#)

Activate Windows
Go to Settings to activate Windows.



phase 4. How to train a machine learning model in a Spark cluster with four EC2 instances running concurrently.

1. When the cluster is ready to accept jobs, you can either use the step button to add steps or submit manually.


```
ssh -i "ec2key.pem" User>>@Public IPv4 DNS>>
```

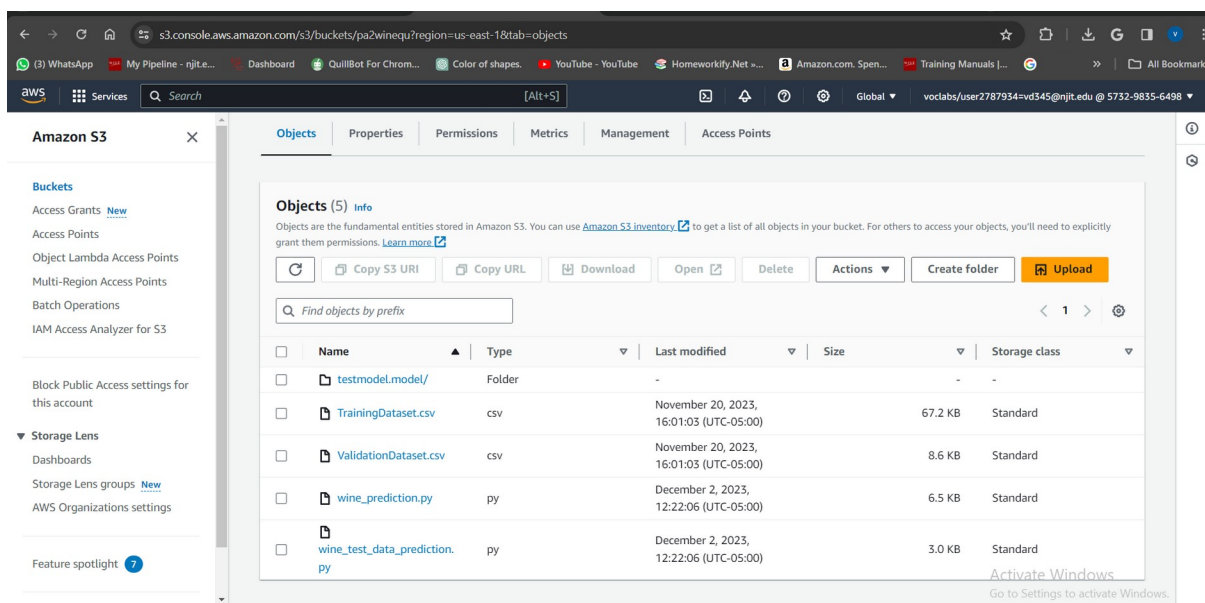
“ sudo su ”

3. Use the following command to submit the job:

```
spark-submit s3://pa2winequ/wine prediction.py
```

```
root@ip-172-31-81-105:/home/ec2-user
23/11/21 21:00:18 INFO SecurityManager: SecurityManager: authentication disabled; ui acls disabled; users with view permissions: Set(root); groups with view permissions: Set(); users with modify permissions: Set(root); groups with modify permissions: Set()
23/11/21 21:00:19 INFO client: Submitting application application_1700599674247_0001 to ResourceManager
23/11/21 21:00:19 INFO YarnClientImpl: Submitted application application_1700599674247_0001
23/11/21 21:00:19 INFO SchedulerExtensionServices: Starting Yarn extension services with app application_1700599674247_0001 and attemptId None
23/11/21 21:00:20 INFO client: Application report for application_1700599674247_0001 (state: ACCEPTED)
23/11/21 21:00:20 INFO client:
  client token: N/A
  diagnostics: AM container is launched, waiting for AM container to register with RM
  ApplicationMaster host: N/A
  ApplicationMaster RPC port: -1
  queue: default
  start time: 1700600419308
  final status: UNDEFINED
  tracking URL: http://ip-172-31-81-105.ec2.internal:20888/proxy/application_1700599674247_0001/
  user: root
23/11/21 21:00:21 INFO client: Application report for application_1700599674247_0001 (state: ACCEPTED)
23/11/21 21:00:22 INFO client: Application report for application_1700599674247_0001 (state: ACCEPTED)
23/11/21 21:00:23 INFO client: Application report for application_1700599674247_0001 (state: ACCEPTED)
23/11/21 21:00:24 INFO YarnClientSchedulerBackend: Add WebUI Filter: org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter, Map(PROXY_HOSTS -> ip-172-31-81-105.ec2.internal, PROXY_URI_BASES -> http://ip-172-31-81-105.ec2.internal:20888/proxy/application_1700599674247_0001/, /proxy/application_1700599674247_0001)
23/11/21 21:00:24 INFO YarnSchedulerBackend$YarnSchedulerEndpoint: ApplicationMaster registered as NettyRpcEndpointRef(spark-client://YarnAM)
23/11/21 21:00:24 INFO client: Application report for application_1700599674247_0001 (state: RUNNING)
23/11/21 21:00:24 INFO client:
  client token: N/A
  diagnostics: N/A
  ApplicationMaster host: 172.31.93.170
  ApplicationMaster RPC port: -1
  queue: default
  start time: 1700600419308
  final status: UNDEFINED
  tracking URL: http://ip-172-31-81-105.ec2.internal:20888/proxy/application_1700599674247_0001/
  user: root
23/11/21 21:00:24 INFO YarnClientSchedulerBackend: Application application_1700599674247_0001 has started running.
23/11/21 21:00:24 INFO client: Successfully started service 'org.apache.spark.network.netty.NettyBlockTransferService' on port 44907.
23/11/21 21:00:24 INFO NettyBlockTransferService: Server created on ip-172-31-81-105.ec2.internal:44907
23/11/21 21:00:24 INFO BlockManager: Using org.apache.spark.storage.RandomBlockReplicationPolicy for block replication policy
23/11/21 21:00:24 INFO BlockManagerMaster: Registering BlockManager BlockManagerId(driver, ip-172-31-81-105.ec2.internal, 44907, None)
23/11/21 21:00:24 INFO BlockManagerMasterEndpoint: Registering block manager ip-172-31-81-105.ec2.internal:44907 with 912.3 MB RAM, BlockManagerId(driver, ip-172-31-81-105.ec2.internal, 44907, None)
23/11/21 21:00:24 INFO BlockManager: external shuffle service port = 7337
23/11/21 21:00:24 INFO BlockManager: Initialized BlockManager: BlockManagerId(driver, ip-172-31-81-105.ec2.internal, 44907, None)
23/11/21 21:00:24 INFO JettyUtils: Adding filter org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter to /metrics/json.
23/11/21 21:00:24 INFO EventLoggingListener: Logging events to hdfs://var/log/spark/apps/application_1700599674247_0001
23/11/21 21:00:24 INFO client: Using initial executors = 50, max of spark.dynamicAllocation.initialExecutors, spark.dynamicAllocation.minExecutors and spark.executor.instances
23/11/21 21:00:24 INFO YarnClientSchedulerBackend: SchedulerBackend is ready for scheduling beginning after reached minRegisteredResourcesRatio: 0.0
23/11/21 21:00:25 INFO SharedState: loading hive config file: file:/etc/spark/conf/hive-site.xml
23/11/21 21:00:25 INFO SharedState: Setting hive metastore warehouse dir to null from the value of spark.sql.warehouse.dir ('hdfs://user/spark/warehouse').
23/11/21 21:00:25 INFO JettyUtils: Adding filter org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter to /SQL.
23/11/21 21:00:25 INFO JettyUtils: Adding filter org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter to /SQL/json.
23/11/21 21:00:25 INFO JettyUtils: Adding filter org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter to /SQL/execution.
23/11/21 21:00:25 INFO JettyUtils: Adding filter org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter to /SQL/execution/json.
23/11/21 21:00:25 INFO JettyUtils: Adding filter org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter to /static/sql.
23/11/21 21:00:25 INFO StateStoreCoordinatorRef: Registered StateStoreCoordinator endpoint
Test Accuracy of wine prediction model = 0.98125
weighted f1 score of wine prediction model = 0.9762234042553191
```

4. You can trace status of this job in EMR UI application logs. Once status is succeeded a test.model will be created in s3 bucket-s3://pa2winequ.



phase 5. Running the Machine learning model without using docker:

- Clone the Repository: Start by cloning this repository onto your local machine.
- Set Up Local Spark Environment: Ensure that you have a local Spark environment ready for running this application. If you haven't set up Spark yet, you can refer to the [official Spark documentation] (<https://spark.apache.org/docs/latest>) for instructions.

- Navigate to the Python File Folder: Go to the 'python file' directory within the cloned repository. Prepare Training Dataset and Place Training Dataset in the 'C:\pa2-cc\data\csv' folder.

phase 6. Run ML model using Docker.

1. Install docker where you want to run this container
2. A public image has been created and posted on DockerHub using cmd –

`docker build -t wine_prediction_app .`

```
PS C:\Users\vijay\pySparkAWSWinePredictionApp-main> docker build -t wine-prediction-app .
[+] Building 124.9s (28/28) FINISHED
=> [internal] load build definition from Dockerfile
=> => transferring dockerfile: 1.27kB
=> [internal] load .dockerignore
=> => transferring context: 2B
=> [internal] load metadata for docker.io/library/centos:7
=> [auth] library/centos:pull token for registry-1.docker.io
=> [ 1/22] FROM docker.io/library/centos:7@sha256:be65f488b7764ad3638f236b7b515b3678369a5124c47b8d32916d6487418ea4
=> => resolve docker.io/library/centos:7@sha256:be65f488b7764ad3638f236b7b515b3678369a5124c47b8d32916d6487418ea4
=> => sha256:eeb6ee3f44bd0b5103bb561b4c16cb82328cfe5809ab675bb17ab3a16c517c9 2.75kB / 2.75kB
=> => sha256:be65f488b7764ad3638f236b7b515b3678369a5124c47b8d32916d6487418ea4 1.20kB / 1.20kB
=> => sha256:dead07b4d8ed7e29e98de0f4504d87e8880d4347859d839686a31da35a3b532f 529B / 529B
=> [internal] load build context
=> => transferring context: 688.19kB
=> [ 2/22] RUN yum -y update && yum -y install python3 python3-dev python3-pip python3-virtualenv java-1.8.0-o
=> [ 3/22] RUN python -V
=> [ 4/22] RUN python3 -V
=> [ 5/22] RUN pip3 install --upgrade pip
=> [ 6/22] RUN pip3 install numpy pandas
=> [ 7/22] RUN pip3 install pandas
=> [ 8/22] RUN wget --no-verbose -O apache-spark.tgz "https://archive.apache.org/dist/spark/spark-3.1.2/spark-3
=> [ 9/22] RUN ln -s /opt/spark-3.1.2-bin-hadoop2.7 /opt/spark
=> [10/22] RUN (echo 'export SPARK_HOME=/opt/spark' >> ~/.bashrc && echo 'export PATH=$SPARK_HOME/bin:$PATH' >>
=> [11/22] RUN mkdir /code
=> [12/22] RUN mkdir /code/data
=> [13/22] RUN mkdir /code/data/csv
=> [14/22] RUN mkdir /code/data/model
=> [15/22] RUN mkdir /code/src
=> [16/22] RUN mkdir /code/data/testdata.model/
=> [17/22] COPY src/wine_test_data_prediction.py /code/src
=> [18/22] COPY data/model/testdata.model/ /code/data/model/testdata.model
=> [19/22] RUN rm /bin/sh && ln -s /bin/bash /bin/sh
=> [20/22] RUN /bin/bash -c "source ~/.bashrc"
=> [21/22] RUN /bin/sh -c "source ~/.bashrc"
=> [22/22] WORKDIR /code
=> => exporting to image
=> => exporting layers
=> => writing image sha256:333701a6bbceaa0c5ebe4409b3fa06683013fb8cfc0685fadbaec84a80769abf
=> => naming to docker.io/library/wine-prediction-app
```

1. You can see the image using cmd
`docker image ls`
2. You can push this in docker hub repository
`docker push vijaya340/wine_prediction_app`
3. A public image has been created and posted on Docker Hub. Use the command
`**Docker pull vijaya340/wine_prediction_app**`
to get the image on your machine.
4. Place your training dataset file in a folder (lets call it directory dirA) , which
you will Count with docker container and run it using below cmd
`docker run -v`
C:/Users/vijay/pySparkAWSWinePredictionApp-main/src:/code/data/csv
winepred TrainingDataset.csv

What's Next?

View a summary of image vulnerabilities and recommendations → [docker scout quickview](#)

```
PS C:\Users\vijay\pySparkAWSWinePredictionApp-main> docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
wine-prediction-app	latest	333701a6bbce	59 seconds ago	1.21GB
ubuntu	latest	e4c58958181a	6 weeks ago	77.8MB
dfordeepika/awssparkwineapp	latest	1d641e9d364e	2 years ago	1.07GB

```
PS C:\Users\vijay\pySparkAWSWinePredictionApp-main> docker run -v C:/Users/vijay/Desktop/pa2:/code/data/csv wine-prediction-app TrainingDataset.csv
```

23/11/21 03:40:25 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties

23/11/21 03:40:29 INFO SparkContext: Running Spark version 3.1.2

23/11/21 03:40:29 INFO ResourceUtils: =====

23/11/21 03:40:29 INFO ResourceUtils: No custom resources configured for spark.driver.

23/11/21 03:40:29 INFO ResourceUtils: =====

23/11/21 03:40:29 INFO SparkContext: Submitted application: tanvi_cs643_wine_prediction

23/11/21 03:40:29 INFO ResourceProfile: Default ResourceProfile created, executor resources: Map(cores -> name: cores, amount: 1, script: , vendor: , memory -> name: memory, amount: 1024, script: , vendor: , offHeap -> name: offHeap, amount: 0, script: , vendor:), task resources: Map(cpus -> name: cpus, amount: 1.0)

23/11/21 03:40:29 INFO ResourceProfile: Limiting resource is cpu

23/11/21 03:40:29 INFO ResourceProfileManager: Added ResourceProfile id: 0

23/11/21 03:40:30 INFO SecurityManager: Changing view acls to: root

23/11/21 03:40:30 INFO SecurityManager: Changing modify acls to: root

23/11/21 03:40:30 INFO SecurityManager: Changing view acls groups to:

23/11/21 03:40:30 INFO SecurityManager: Changing modify acls groups to:

23/11/21 03:40:30 INFO SecurityManager: SecurityManager: authentication disabled; ui acls disabled; users with view permissions: Set(root); groups with view permissions: Set(); users with modify permissions: Set(root); groups with modify permissions: Set()

23/11/21 03:40:30 INFO Utils: Successfully started service 'sparkDriver' on port 36373.

23/11/21 03:40:30 INFO SparkEnv: Registering MapOutputTracker

23/11/21 03:40:30 INFO SparkEnv: Registering BlockManagerMaster

23/11/21 03:40:30 INFO BlockManagerMasterEndpoint: Using org.apache.spark.storage.DefaultTopologyMapper for getting topology information

23/11/21 03:40:30 INFO BlockManagerMasterEndpoint: BlockManagerMasterEndpoint up

23/11/21 03:40:30 INFO SparkEnv: Registering BlockManagerMasterHeartbeat

23/11/21 03:40:30 INFO DiskBlockManager: Created local directory at /tmp/blockmgr-a622e406-16d8-47ad-9c67-56b322f7f45f

23/11/21 03:40:30 INFO MemoryStore: MemoryStore started with capacity 366.3 MiB

23/11/21 03:40:30 INFO SparkEnv: Registering OutputCommitCoordinator

23/11/21 03:40:31 INFO Utils: Successfully started service 'SparkUI' on port 4040.

23/11/21 03:40:31 INFO SparkUI: Bound SparkUI to 0.0.0.0, and started at http://0369ac72f56c:4040

23/11/21 03:40:32 INFO Executor: Starting executor ID driver on host 0369ac72f56c

23/11/21 03:40:32 INFO Utils: Successfully started service 'org.apache.spark.network.netty.NettyBlockTransferService' on port 32979.

23/11/21 03:40:32 INFO NettyBlockTransferService: Server created on 0369ac72f56c:32979

Activate Windows

Go to Settings to activate Windows.

phase 7. Conclusion

After running the models with docker the quality of wine predicted to be

	With Docker
Accuracy	98.358
F1 Score	97.766

From the above table we concluded with docker the accuracy of the model is yielding accuracy is 98.358